

GENERATIVE ADVERSARIAL NETWORKS (GANs) - SPECIALIZATION

COURSE 1 - BUILD BASIC GENERATIVE ADVERSARIAL NETWORKS

GENERATIVE MODELS

- DISCRIMINATIVE MODELS: FROM A FEATURE X , DISCRIMINATE TO WHICH CLASS Y THIS FEATURE BELONGS.

$$\text{FEATURES}(X) \rightarrow \text{CLASS}(Y) \quad P(Y|X)$$

- GENERATIVE MODELS: FROM A Y LABEL AND A RANDOM NOISE VECTOR, IT GENERATES A FEATURE X .

$$\text{NOISE}(\epsilon), \text{CLASS}(y) \rightarrow \text{FEATURES}(X) \quad P(X, Y)$$

↳ VARIATIONAL AUTOENCODERS ENCODER \rightarrow LATENT SPACE \rightarrow DECODER AFTER TRAINING, THE ENCODER IS
LATENT SPACE \rightarrow DECODER \rightarrow RESULT NO LONGER NEEDED.

↳ VARIES ACCESS TO DATA DISTRIBUTION TO PRODUCE DIFFERENT RESULTS.

↳ GENERATIVE ADVERSARIAL NETWORKS GENERATOR DISCRIMINATOR

RANDOM NOISE \rightarrow GENERATOR \rightarrow RESULT COMPETE WITH EACH OTHER.

↳ THE DISCRIMINATOR IS ONLY USEFUL DURING TRAINING.

↳ GENERATIVE MODELS LEARN TO PRODUCE REALISTIC EXAMPLES AND DISCRIMINATIVE MODELS DISTINGUISH
BETWEEN CLASSES.

REAL LIFE GANS

↳ GANs PERFORMANCE IS RAPIDLY IMPROVING.

- FACE GENERATION PROGRESS: BLACK AND WHITE (2014) \rightarrow COLORED PHOTO-REALISTIC (2018)

- MIMICS THE DISTRIBUTION OF THE TRAINING DATA. (STYLE GAN)

- GANs FOR IMAGE TRANSLATION: FROM ONE DOMAIN TO ANOTHER. (CYCLED GAN)

- FROM A DRAWING, GENERATE A REALISTIC IMAGE. (GAUGAN)

- ANIMATE FACE IMAGES, MAKING THEM HAVE REACTIONS.

- GANs FOR 3D OBJECTS. (3D-GAN)

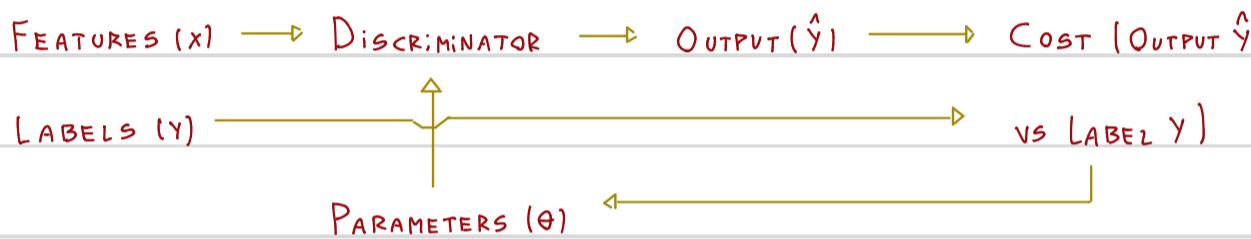
- COMPANIES USING GANs: ADOBE (NEXT-GEN PHOTOSHOP), GOOGLE (TEXT GENERATION), IBM (DATA AUGMENTATION),
SNAPCHAT / TIK TOK (IMAGE FILTERS) AND DISNEY (SUPER-RESOLUTION).

INTUITION BEHIND GANS

- GENERATOR LEARNS TO MAKE FAKES THAT LOOK REAL. → DOESN'T KNOW HOW IT SHOULD LOOK.
- DISCRIMINATOR LEARNS TO DISTINGUISH REAL FROM FAKE.
- THE GENERATOR'S GOAL IS TO FOOL THE DISCRIMINATOR.
- THE DISCRIMINATOR'S GOAL IS TO DISTINGUISH BETWEEN REAL AND FAKE.
- THEY LEARN FROM THE COMPETITION WITH EACH OTHER.
- AT THE END, FAKES LOOK REAL.

DISCRIMINATOR

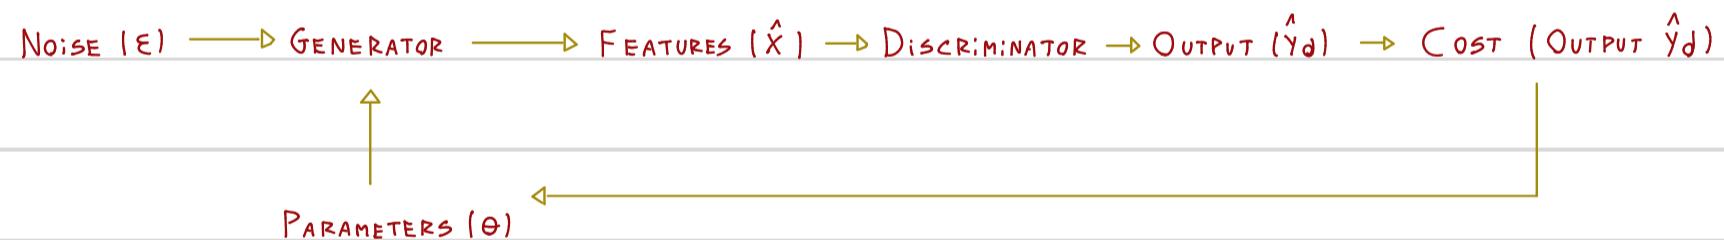
- CLASSIFIERS: DISTINGUISH BETWEEN DIFFERENT CLASSES.



- IT LEARNS THE PROBABILITY OF CLASS Y (REAL OR FAKE) GIVEN FEATURES X.

- THE PROBABILITIES ARE THE FEEDBACK FOR THE GENERATOR.

- GENERATOR: GENERATES EXAMPLES OF THE CLASS.



- SAMPLING NOISE VECTOR: DIFFERENT OUTPUTS AT EVERY RUN.

- GENERATES IMAGES THAT APPROXIMATE WITH THE DATA DISTRIBUTION OF THE TRAINING SET.

- THE GENERATOR PRODUCES FAKE DATA TAKES AS INPUT NOISE (RANDOM FEATURES).

- IT LEARNS THE PROBABILITY OF FEATURES X.

BCE Cost Function (Binary Cross Entropy)

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1-y^{(i)}) \log (1-h(x^{(i)}, \theta))]$$

↳ ENSURES THAT THE COST IS ALWAYS GREATER OR EQUAL TO 0
↳ PREDICTION
↳ FEATURES
↳ LABEL
↳ PARAMETERS
↳ AVERAGE LOSS OF THE WHOLE BATCH

$y^{(i)}$	$h(x^{(i)}, \theta)$	$y^{(i)} \log h(x^{(i)}, \theta)$	$y^{(i)}$	$h(x^{(i)}, \theta)$	$(1-y^{(i)}) \log (1-h(x^{(i)}, \theta))$
0	ANY	0	1	ANY	0
1	0,99	~0	0	0,01	~0
1	~0	-∞	0	~1	-∞

↳ RELEVANT WHEN THE LABEL IS 1

↳ RELEVANT WHEN THE LABEL IS 0

CLOSE TO ZERO WHEN THE LABEL AND THE PREDICTION ARE SIMILAR.

APPROACHES INFINITY WHEN THE LABEL AND THE PREDICTION ARE DIFFERENT.

PUTTING IT ALL TOGETHER

DISCRIMINATOR LEARN GETTING FEEDBACK ON IF ITS CLASSIFICATION WAS CORRECT.

GENERATOR LEARN USING FEEDBACK FROM THE DISCRIMINATOR.

SUPERIOR DISCRIMINATOR → FAKE AS 100% FAKE → NO WAY TO IMPROVE.

↳ THE TWO MODELS SHOULD ALWAYS BE AT A SIMILAR "SKILL" LEVEL.

NOISE (ϵ) → GENERATOR → \hat{x} → DISCRIMINATOR → OUTPUT (\hat{y})



INTRO TO PyTorch

PyTorch: IMPERATIVE, COMPUTATIONS ON THE GO;

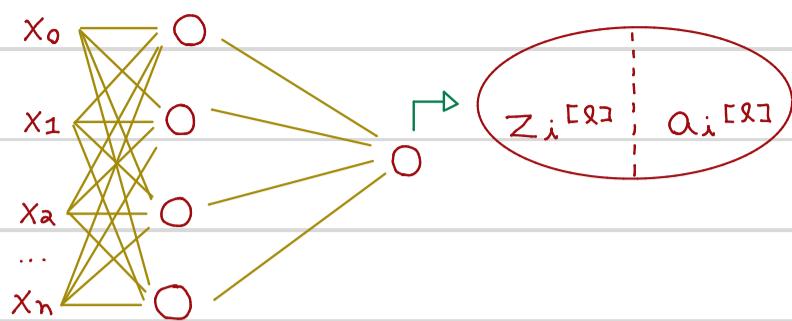
DYNAMIC COMPUTATIONAL GRAPHS;

TensorFlow: SYMBOLIC, FIRST DEFINE AND THEN COMPILE;

STATIC COMPUTATIONAL GRAPHS;

TensorFlow > 2.0 moves toward PyTorch by including EAGER EXECUTION.

• ACTIVATIONS (BASIC PROPERTIES)



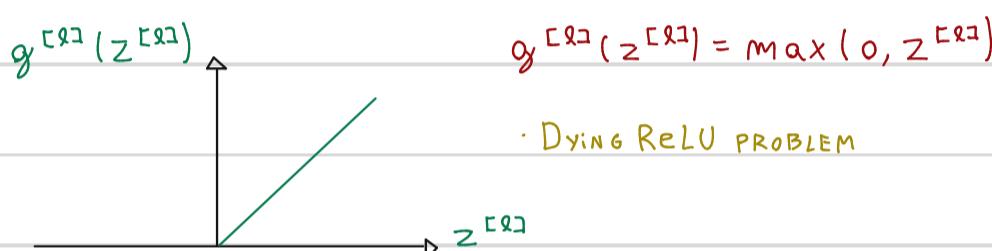
$$z_i = \sum_{i=0}^n w_i x_i$$

$a_i = g(z_i) \rightarrow$ DIFFERENTIABLE NON-LINEAR FUNCTION.

- DIFFERENTIABLE FOR BACKPROPAGATION.
- NON-LINEAR TO COMPUTE COMPLEX FEATURES.

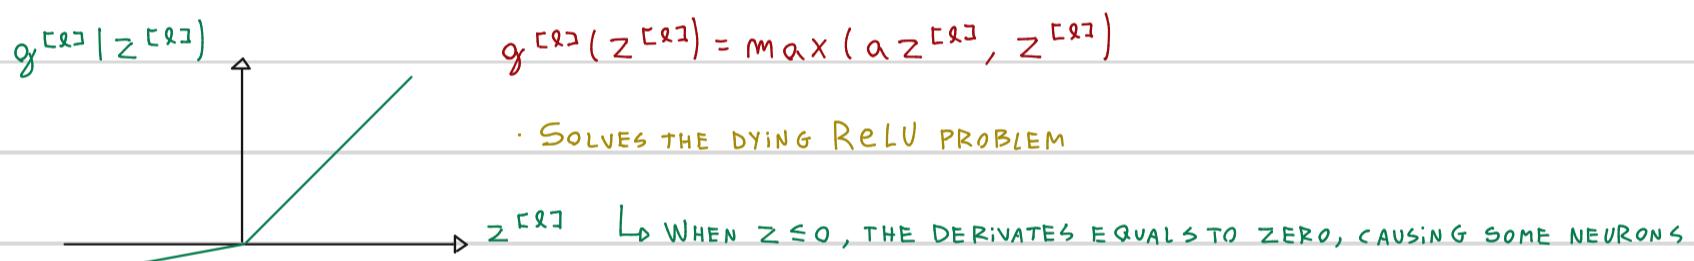
• COMMON ACTIVATION FUNCTIONS

• ACTIVATIONS: ReLU = RECTIFIED LINEAR UNIT



• Dying ReLU PROBLEM

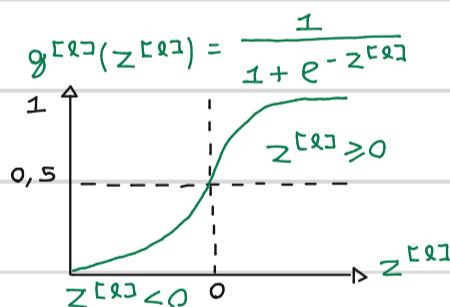
• ACTIVATIONS: LEAKY ReLU



• SOLVES THE DYING ReLU PROBLEM

• WHEN $z \leq 0$, THE DERIVATIVE EQUALS TO ZERO, CAUSING SOME NEURONS TO GET STUCK AND STOP LEARNING.

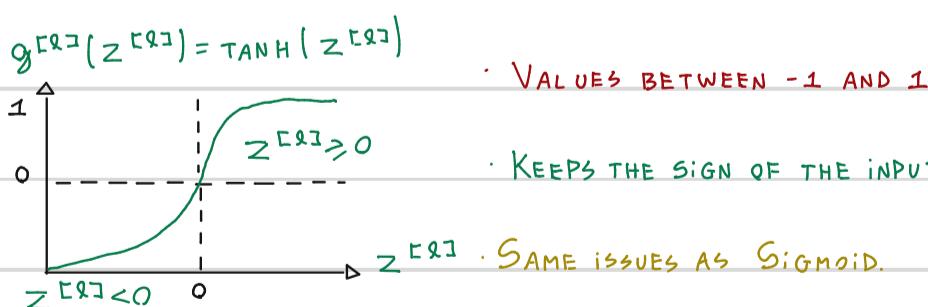
• ACTIVATIONS: SIGMOID



• VALUES BETWEEN 0 AND 1.

• VANISHING GRADIENT AND SATURATION PROBLEMS.

• ACTIVATIONS: TANH



• VALUES BETWEEN -1 AND 1.

• KEEPS THE SIGN OF THE INPUT.

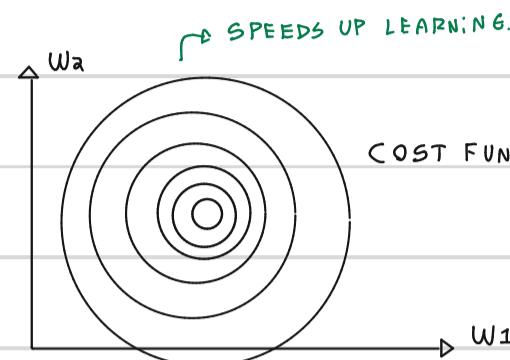
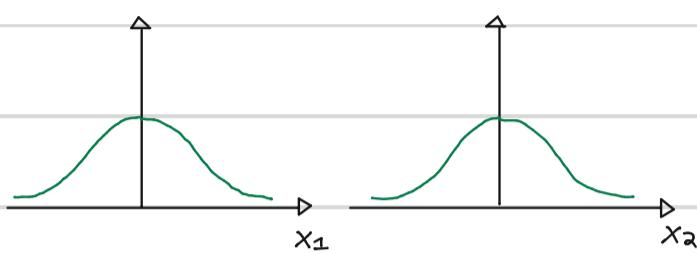
• SAME ISSUES AS Sigmoid.

• VANISHING GRADIENT PROBLEM: ACTIVATION FUNCTIONS LIKE THE SIGMOID HAVE A DERIVATIVE THAT IS BOUNDED BETWEEN 0 AND 1, MEANING THAT THE GRADIENT DECREASES EXPONENTIALLY AS INFORMATION IS PROPAGATED THROUGH THE VARIOUS LAYERS OF THE NEURAL NETWORKS. AS A RESULT, THE DEEPER LAYERS OF THE NEURAL NETWORKS CAN RECEIVE VERY SMALL GRADIENTS, MAKING IT DIFFICULT OR IMPOSSIBLE TO UPDATE THEIR WEIGHTS IN A MEANINGFUL WAY.

- SATURATION GRADIENT PROBLEM: OCCURS WHEN THE ACTIVATION FUNCTION OF A LAYER OF THE NEURAL NETWORK "SATURATES" OR APPROACHES ITS UPPER OR LOWER LIMIT, IN SUCH A WAY THAT THE GRADIENT BECOMES TOO SMALL TO UPDATE THE WEIGHTS OF THAT LAYER EFFECTIVELY.

- BATCH NORMALIZATION (EXPLAINED)

NORMALIZATION AND ITS EFFECTS



- TRAINING DATA USES BATCH STATS.
- TEST DATA USES TRAINING STATS.

↳ REDUCTION OF COVARIATE SHIFT. → COVARIATE SHIFT OCCURS WHEN TRAINING AND TEST DATA HAVE DIFFERENT DISTRIBUTIONS, WHICH CAN LEAD TO INACCURATE OR UNEXPECTED MODEL RESULTS.

↳ SMOOTHES THE COST FUNCTION.

- BATCH NORMALIZATION (PROCEDURE)

$$\hat{z}_i^{[q]} = \frac{z_i^{[q]} - \bar{z}_i^{[q]}}{\sqrt{\sigma_{z_i^{[q]}}^2 + \epsilon}}$$

BATCH MEAN OF $z_i^{[q]}$

BATCH STD OF $z_i^{[q]}$

FOR EVERY EXAMPLE IN THE BATCH TRAINING

$$y_i^{[q]} = \gamma \hat{z}_i^{[q]} + \beta$$

SHIFT FACTOR

SCALE FACTOR → LEARNABLE PARAMETERS TO GET THE OPTIMAL DIST.

$$\hat{z}_i^{[q]} = \frac{z_i^{[q]} - E(z_i^{[q]})}{\sqrt{Var(z_i^{[q]}) + \epsilon}}$$

RUNNING MEAN AND RUNNING STD FROM TRAINING.

FRAMEWORKS LIKE TENSORFLOW AND PYTORCH KEEP TRACK OF THEM.

FOR EVERY EXAMPLE IN THE BATCH TEST

- REVIEW OF CONVOLUTIONS

- CONVOLUTIONS ARE USEFUL LAYERS FOR PROCESSING IMAGES.

- THEY SCAN THE IMAGE TO DETECT USEFUL FEATURES.

- JUST ELEMENT-WISE PRODUCTS AND SUMS.

- PADDING AND STRIDE

- STRIDE DETERMINES HOW THE FILTER SCANS THE IMAGE.

- PADDING LIKE A FRAME ON THE IMAGE.

- PADDING GIVES SIMILAR IMPORTANCE TO THE EDGES AND THE CENTER.

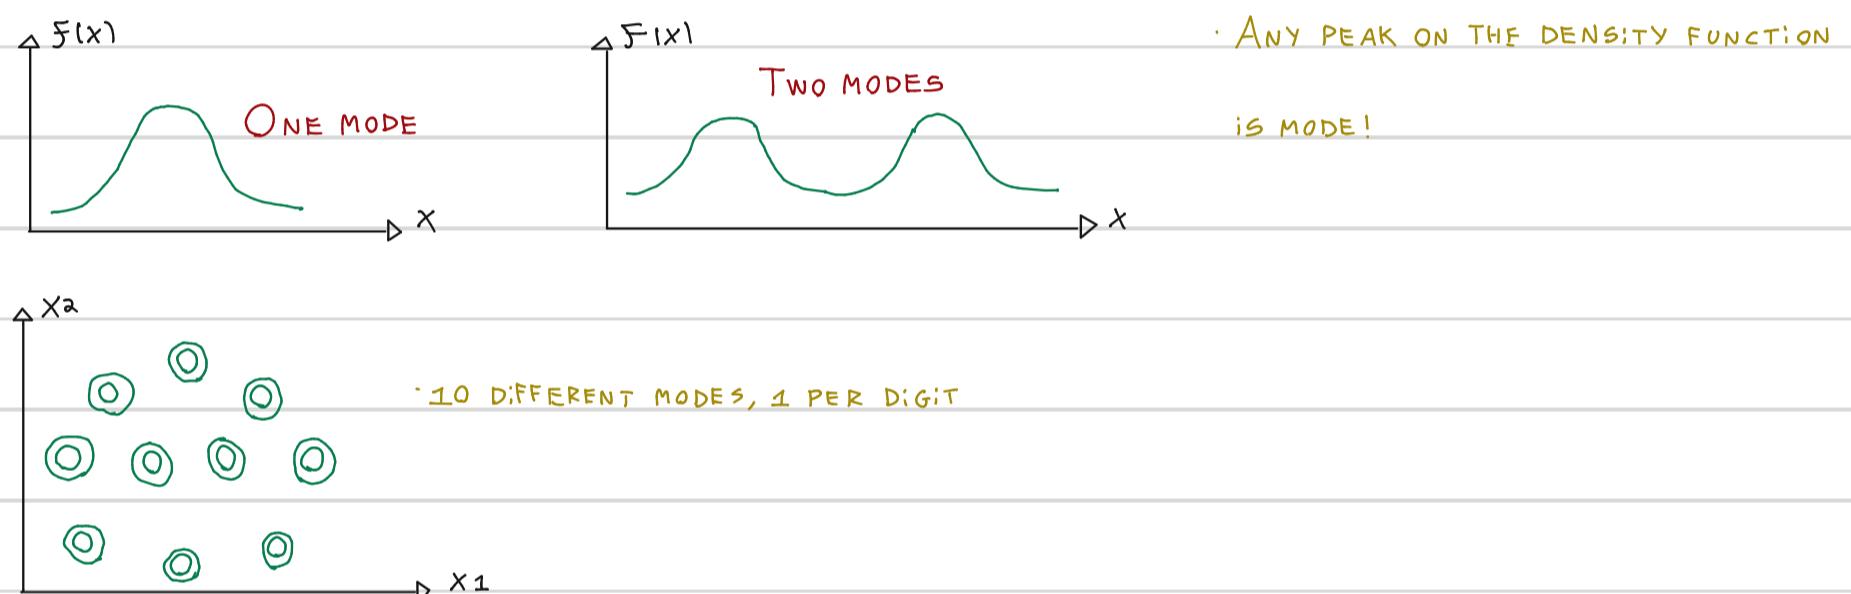
• POOLING AND UPSAMPLING

- TYPES OF POOLING: MAX POOLING, AVERAGE POOLING, MIN POOLING, ...
- TYPES OF UPSAMPLING: NEAREST NEIGHBORS, LINEAR INTERPOLATION, BI-LINEAR INTERPOLATION, ...
- NO LEARNABLE PARAMETERS.

• TRANSPOSED CONVOLUTIONS

- TRANSPOSED CONVOLUTIONS UPSAMPLE.
- THEY HAVE LEARNABLE PARAMETERS.
- PROBLEM: RESULTS HAVE A CHECKERBOARD PATTERN.

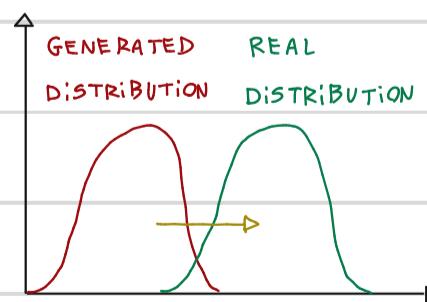
• MODE COLLAPSE



- MODES ARE PEAKS IN THE DISTRIBUTION OF FEATURES.
- TYPICAL WITH REAL-WORLD DATASETS.
- MODE COLLAPSE HAPPENS WHEN THE GENERATOR GETS STUCK IN ONE MODE (THE GENERATOR GETS STUCK IN A LOCAL MINIMA). THE DISCRIMINATOR WILL EVENTUALLY LEARN TO DIFFERENTIATE THE GENERATOR'S FAKES WHEN THIS HAPPENS AND OUTSKILL IT, ENDING THE MODEL'S LEARNING.

• PROBLEM WITH BCE LOSS

- GENERATOR: MAXIMIZE COST.
- DISCRIMINATOR: MINIMIZE COST.
- DISCRIMINATOR: SINGLE OUTPUT, EASIER TO TRAIN THAN THE GENERATOR. \rightarrow CRITICIZING IS MORE STRAIGHTFORWARD.
- GENERATOR: COMPLEX OUTPUT, DIFFICULT TO TRAIN.



• MAKE THE GENERATED AND REAL DISTRIBUTIONS LOOK SIMILAR.

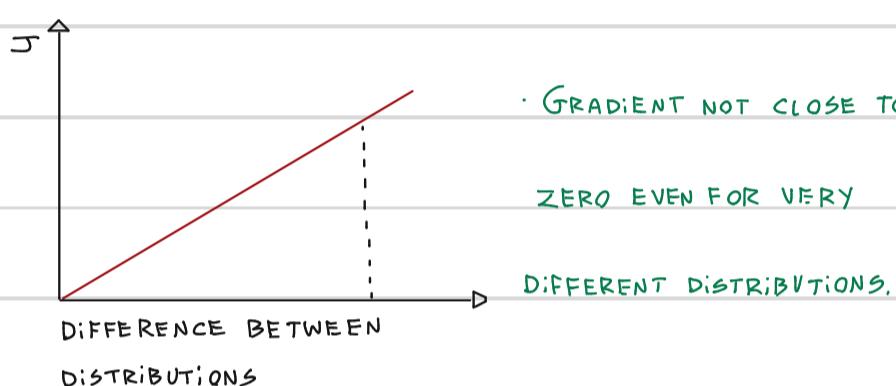
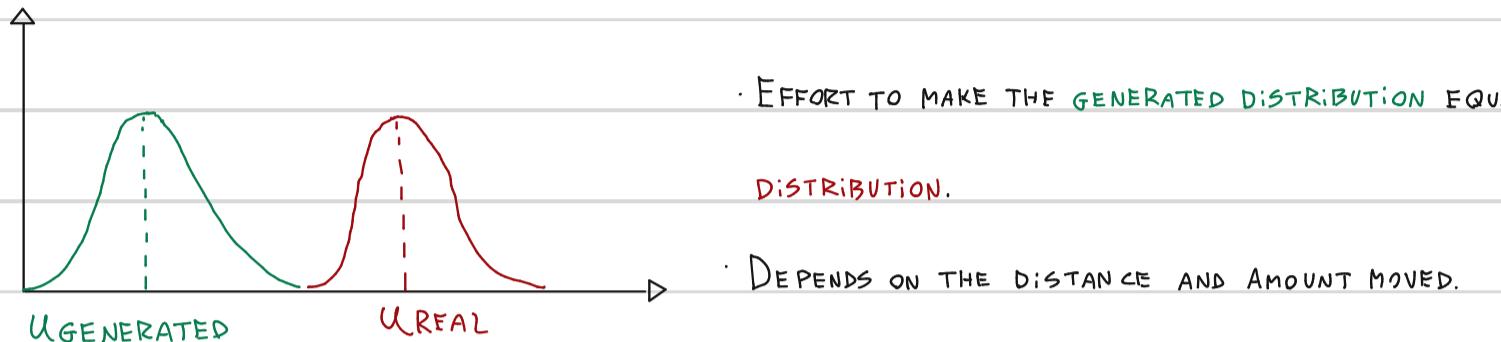
\hookrightarrow OFTEN, THE DISCRIMINATOR GETS BETTER THAN THE GENERATOR.

- GANs TRY TO MAKE THE REAL AND GENERATED DISTRIBUTIONS LOOK SIMILAR.

- WHEN THE DISCRIMINATOR IMPROVES TOO MUCH, THE FUNCTION APPROXIMATED BY BCE LOSS WILL CONTAIN FLAT REGIONS.

- FLAT REGIONS ON THE COST FUNCTION = VANISHING GRADIENTS.

EARTH MOVER'S DISTANCE



- EARTH MOVER'S DISTANCE (EMD) IS A FUNCTION OF AMOUNT AND DISTANCE.

- DOESN'T HAVE FLAT REGIONS WHEN THE DISTRIBUTIONS ARE VERY DIFFERENT.

- APPROXIMATING EMD SOLVES THE PROBLEMS ASSOCIATED WITH BCE.

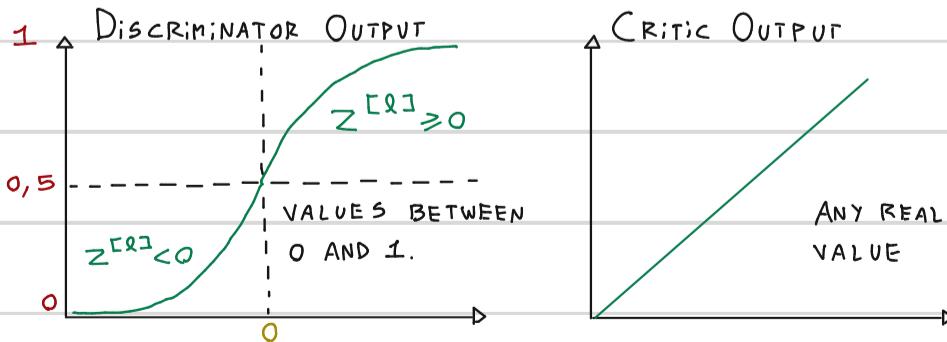
WASSERSTEIN LOSS

- WITH BCE LOSS: DISCRIMINATOR \rightarrow MINIMIZE COST, GENERATOR \rightarrow MAXIMIZE COST.

- W-LOSS APPROXIMATES THE EARTH MOVER'S DISTANCE. $\min_{\mathcal{G}} \max_{\mathcal{C}} \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$

GENERATOR \rightarrow MINIMIZE THE DISTANCE.

CITIC \rightarrow MAXIMIZE THE DISTANCE.



- BCE Loss \rightarrow DISCRIMINATOR OUTPUTS BETWEEN 0 AND 1

$$- [\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z)))]$$

- W-LOSS \rightarrow CRITIC OUTPUTS ANY NUMBER

$$\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

- W-LOSS HELPS WITH MODE COLLAPSE AND VANISHING GRADIENT PROBLEMS.

CONDITION ON WASSERSTEIN CRITIC

$$\min_{g \in \mathcal{C}} \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

NEEDS TO BE 1-LIPSCHITZ CONTINUOUS

THE NORM OF THE GRADIENT SHOULD BE AT MOST 1 FOR EVERY POINT.

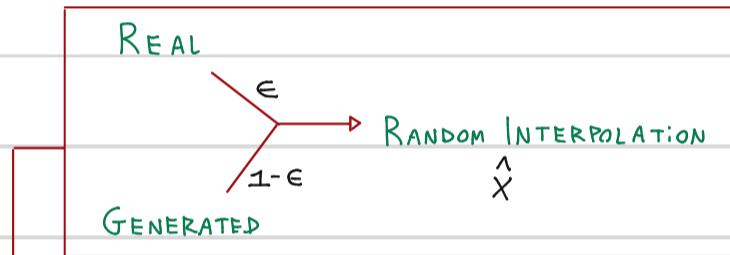
NEEDED FOR TRAINING STABLE NEURAL NETWORKS WITH W-LOSS. → THIS CONDITION ENSURES THAT W-LOSS IS VALIDLY APPROXIMATING EARTH MOVER'S DISTANCE.

1-LIPSCHITZ CONTINUITY ENFORCEMENT

WEIGHT CLIPPING FORCES THE WEIGHTS OF THE CRITIC TO A FIXED INTERVAL.

GRADIENT DESCENT TO UPDATE WEIGHTS → CLIP THE CRITIC'S WEIGHTS. → LIMITS THE LEARNING ABILITY OF THE CRITIC.

GRADIENT PENALTY: $\min_{g \in \mathcal{C}} \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$ → REGULARIZATION OF THE CRITIC'S GRADIENT.



$(\|\nabla c(\hat{x})\|_2 - 1)^2$ → REGULARIZATION TERM
 $\epsilon x + (1-\epsilon)g(z)$ → INTERPOLATION

PUTTING IT ALL TOGETHER: $\min_{g \in \mathcal{C}} \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$

(a) MAKES THE GAN LESS PRONE TO MODE COLLAPSE AND VANISHING GRADIENT.

(b) TRIES TO MAKE THE CRITIC BE 1-L CONTINUOUS, FOR THE LOSS FUNCTION TO BE CONTINUOUS AND DIFFERENTIABLE.

WEIGHT CLIPPING AND GRADIENT PENALTY ARE WAYS TO ENFORCE 1-L CONTINUITY.

→ USING THE GRADIENTS ON AN INTERMEDIATE IMAGE WITH RESPECT TO THE CRITIC IS AN APPROXIMATION FOR ENFORCING THE GRADIENT

NORM TO BE 1 ALMOST EVERYWHERE. → SINCE CHECKING THE CRITIC'S GRADIENT AT EACH POINT OF THE FEATURE SPACE IS VIRTUALLY IMPOSSIBLE, YOU CAN APPROXIMATE THIS BY USING INTERPOLATED IMAGES.

CONDITIONAL GENERATION: INTUITION

- UNCONDITIONAL GENERATION: YOU GET OUTPUTS FROM A RANDOM CLASS. INPUT → GENERATOR → GENERATED EXAMPLE
- CONDITIONAL GENERATION: YOU GET WHAT YOU ASK FOR. INPUT → GENERATOR → GENERATED EXAMPLE
CLASS
- CONDITIONAL:
 - EXAMPLES FROM THE CLASSES YOU WANT.
 - TRAINING DATASET NEEDS TO BE LABELED.
- UNCONDITIONAL:
 - EXAMPLES FROM RANDOM CLASSES.
 - TRAINING DATASET DOESN'T NEED TO BE LABELED.

CONDITIONAL GENERATION: INPUTS

- NOISE VECTOR + CLASS (ONE-HOT) VECTOR → GENERATOR → GENERATED EXAMPLE

↳ CONTROL IN THE GENERATION;
 ↳ RANDOMNESS IN THE GENERATION;

- IMAGE + LABEL → DISCRIMINATOR → OUTPUT PREVISION

DISCRIMINATOR INPUT: IMAGE + CLASS (ONE-HOT) → MATRICES FULL OF ZEROS, IF NOT A CLASS, AND MATRIX FULL OF ONES IF IS A CLASS;
 ↳ 3 CHANNELS (OR 1 IF GRayscale);

CONTROLLABLE GENERATION

- CHANGE SPECIFIC FEATURES OF THE OUTPUT.
- TWEAK THE INPUT NOISE VECTOR TO GET DIFFERENT FEATURES ON THE OUTPUT.

NOISE VECTOR (Z) → GENERATOR → CONTROLLED OUTPUT

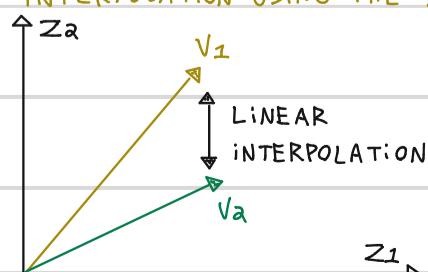
- CONTROLLABLE:

- CONDITIONAL:

- EXAMPLES WITH THE FEATURES THAT YOU WANT.
- TRAINING DATASET DOESN'T NEED TO BE LABELED.
- MANIPULATE THE Z VECTOR INPUT.
- EXAMPLES FROM THE CLASSES YOU WANT.
- TRAINING DATASET NEEDS TO BE LABELED.
- APPEND A CLASS VECTOR TO THE INPUT.

VECTOR ALGEBRA IN THE Z -SPACE

- INTERPOLATION USING THE Z -SPACE: HOW AN IMAGE MORPHS INTO ANOTHER.



Z -SPACE WITH NOISE VECTORS

$g(V_1)$ ← → $g(V_a)$
INTERMEDIATE IMAGES USING THE Z -SPACE

- TO CONTROL OUTPUT FEATURES, YOU NEED TO

FIND DIRECTIONS IN THE Z -SPACE.

- TO MODIFY YOUR OUTPUT, YOU MOVE AROUND IN THE Z -SPACE.

CHALLENGES WITH CONTROLLABLE GENERATION

FEATURE CORRELATION

- UNCORRELATED FEATURES

WOMAN \rightarrow ADD BEARD \rightarrow WOMAN WITH BEARD

- CORRELATED FEATURES

WOMAN \rightarrow ADD BEARD \rightarrow MAN WITH BEARD



MAKE MORE MASCULINE

Z-SPACE ENTANGLEMENT: IT IS NOT POSSIBLE TO CONTROL SINGLE OUTPUT FEATURES.

- WHEN TRYING TO CONTROL ONE FEATURE, OTHERS THAT ARE CORRELATED CHANGE.

- Z-SPACE ENTANGLEMENT MAKES CONTROLLABILITY DIFFICULT, IF NOT IMPOSSIBLE.

- ENTANGLEMENT HAPPENS WHEN Z DOES NOT HAVE ENOUGH DIMENSIONS.

CLASSIFIER GRADIENTS



DIRECTION TO MAKE THE NOISE VECTOR UNTIL THE .92 SUNGLASSES

FEATURE EMERGES.

- CLASSIFIERS CAN BE USED TO FIND DIRECTIONS IN THE Z-SPACE.

- TO FIND DIRECTIONS, THE UPDATES ARE DONE JUST TO THE NOISE VECTOR.

DISENTANGLEMENT

DISENTANGLING Z-SPACES

LET YOU CONTROL INDIVIDUAL FEATURES BY CORRESPONDING Z VALUES DIRECTLY TO THEM.

- THERE ARE SUPERVISED AND UNSUPERVISED METHODS TO ACHIEVE DISENTANGLEMENT.

DISENTANGLING Z-SPACES:

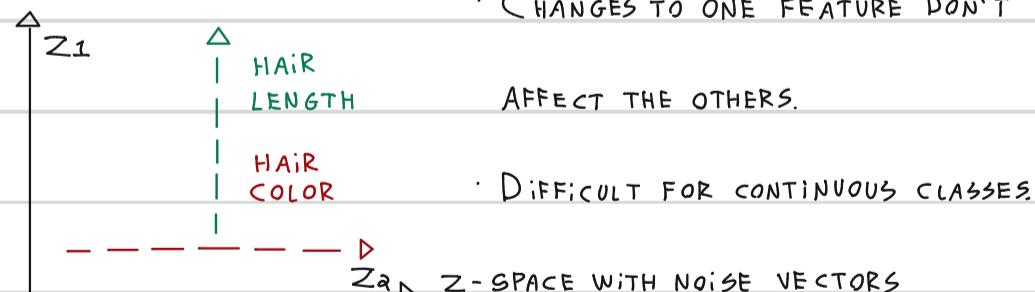
$z_1 \quad z_2$

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

HAIR COLOR

HAIR LENGTH \rightarrow LATENT FACTORS OF VARIATION.



$$L_{\text{NEW}} = L_{\text{ORIGINAL}} + \text{reg}_d$$

ORIGINAL LOSS

REGULARIZATION

CAN BE ANY LOSS FUNCTION

(e.g. BCE, W-LOSS).